

Figure 1 - Overall ENER Pipeline




Figure 1 illustrates an end-to-end encrypted neural embedding relay pipeline. Neural telemetry is acquired locally, converted into a compressed latent embedding, encrypted before relay, and processed by an encrypted inference engine without exposing raw neural telemetry to the remote processor.

Figure 2 - Split-Device Architecture




Figure 2 illustrates a split-device architecture in which neural acquisition and edge encoding occur locally, encrypted embeddings are transmitted to an encrypted compute server, and the result is decrypted or displayed only by an authorized client application.

Figure 3 - Adaptive Latent Compression Controller




Figure 3 illustrates an adaptive compression controller that converts signal-quality, bandwidth, encryption-budget, entropy, privacy-budget, and task-type inputs into a compression policy controlling latent dimensionality, quantization, sparsity, padding, ciphertext packing, and routing.

Figure 4 - Reconstruction-Resistant Encoder Training




Figure 4 illustrates training or configuring an encoder to preserve task utility while reducing reconstruction of raw neural signals. An adversarial decoder attempts reconstruction from the latent embedding, and the encoder is penalized when reconstruction or identity leakage succeeds.

Figure 5 - Encrypted Cognitive-State Classifier




Figure 5 illustrates an encrypted classifier in which encrypted latent vectors are processed through encrypted model operations to produce encrypted scores or classification results. Only an authorized endpoint decrypts or consumes the result.